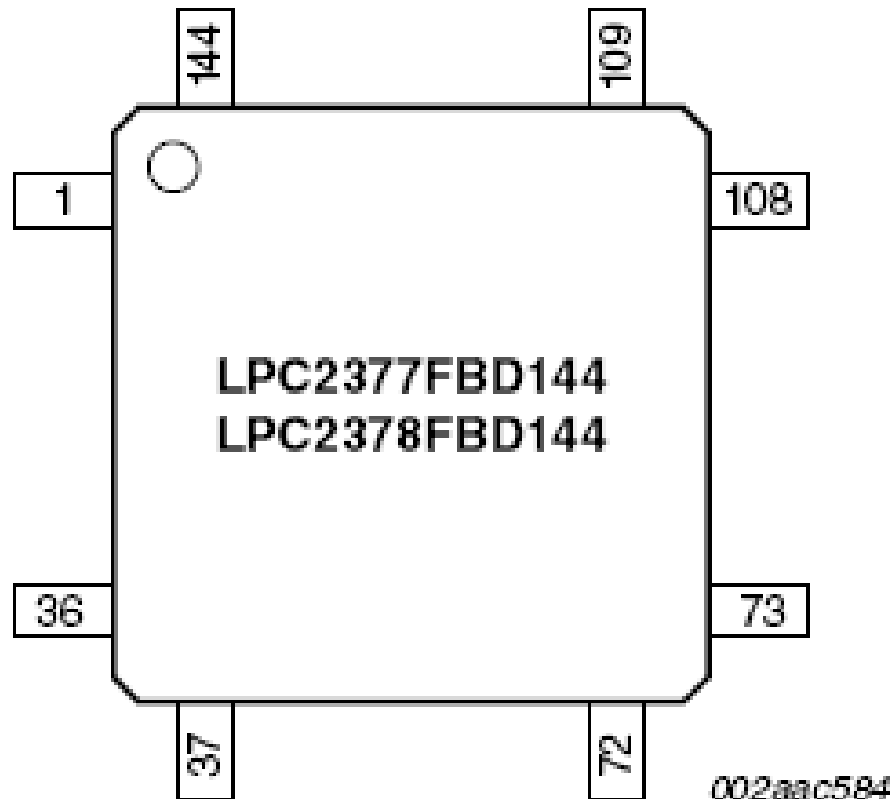


Embedded Systems and Microprocessor Systems

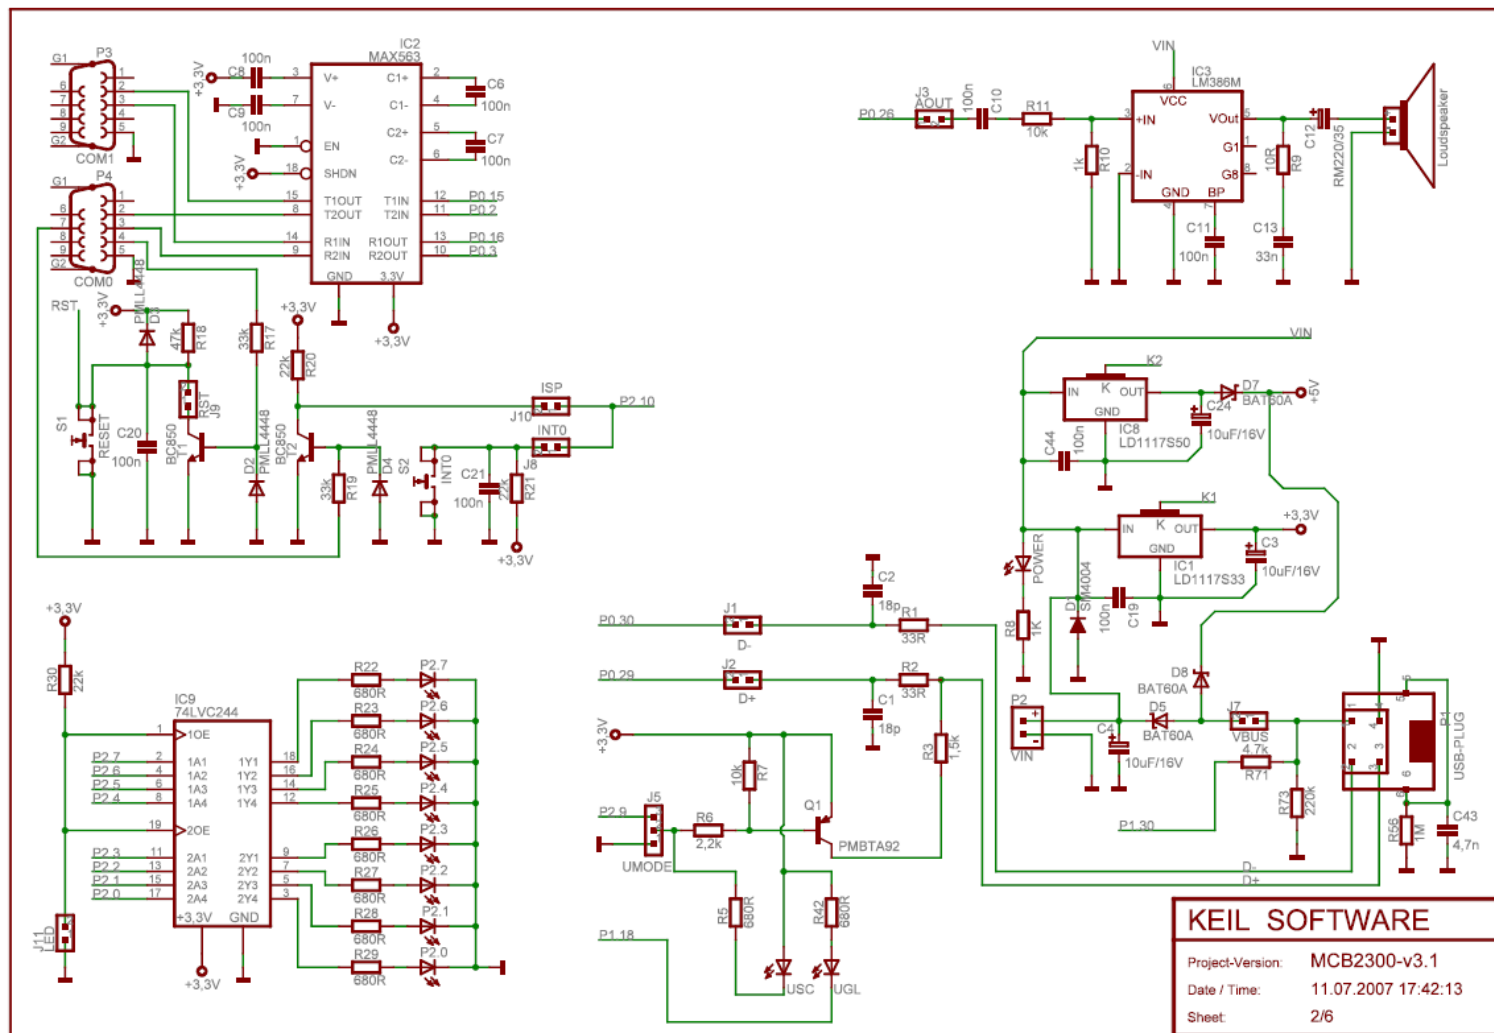
review

- ✓ Organization/Architecture
- ✓ Mnemonics
- ✓ Flags
- ✓ Representations
- ✓ Subroutines
- ✓ Interrupt
- ✓ Pipeline
- ✓ Encoding
- ✓ Memory
- ✓ Processor Modes

- ✓ **Organization/Architecture**
- ✓ Mnemonics
- ✓ Flags
- ✓ Representations
- ✓ Subroutines
- ✓ Interrupt
- ✓ Pipeline
- ✓ Encoding
- ✓ Memory
- ✓ Processor Modes







ARM7TDMI microprocessor




- 74,209 transistors
- 32 bit address bus
- 0.13 μm gate length
- upto 133 MHz clock
- 0.26 mm^2 die area
- max. power 8 mW

T: 16-bit **T**humb code
D: on-chip **D**ebug support
M: an enhanced **M**ultiplier
I: Embedded **I**CE hardware to give on-chip breakpoint and watchpoint support

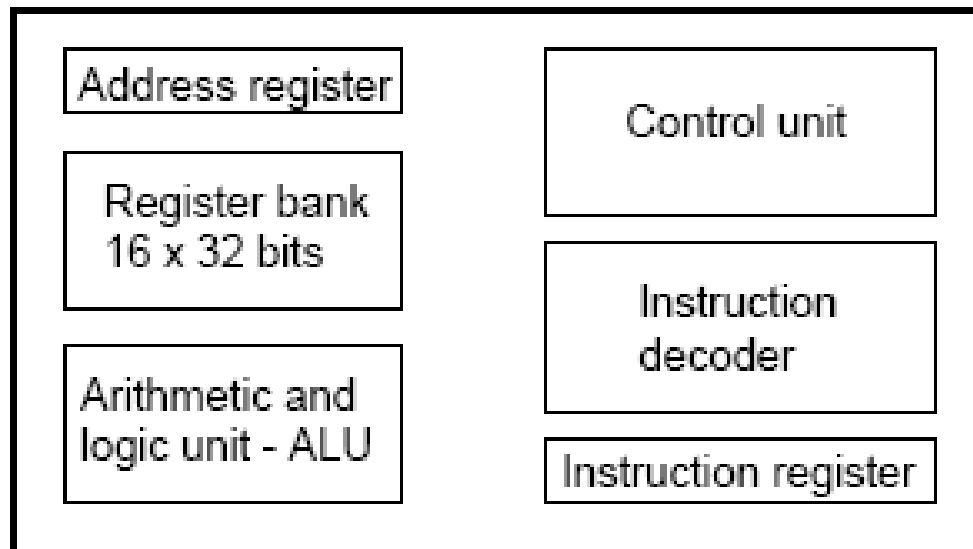
A Third Bus

- In addition to the data bus and control bus there is a third bus called the **address bus**.
- The address bus is used by the CPU to determine which location in memory is sending or receiving data.



The ARM7 core

- The basic building blocks of the ARM7 core are:

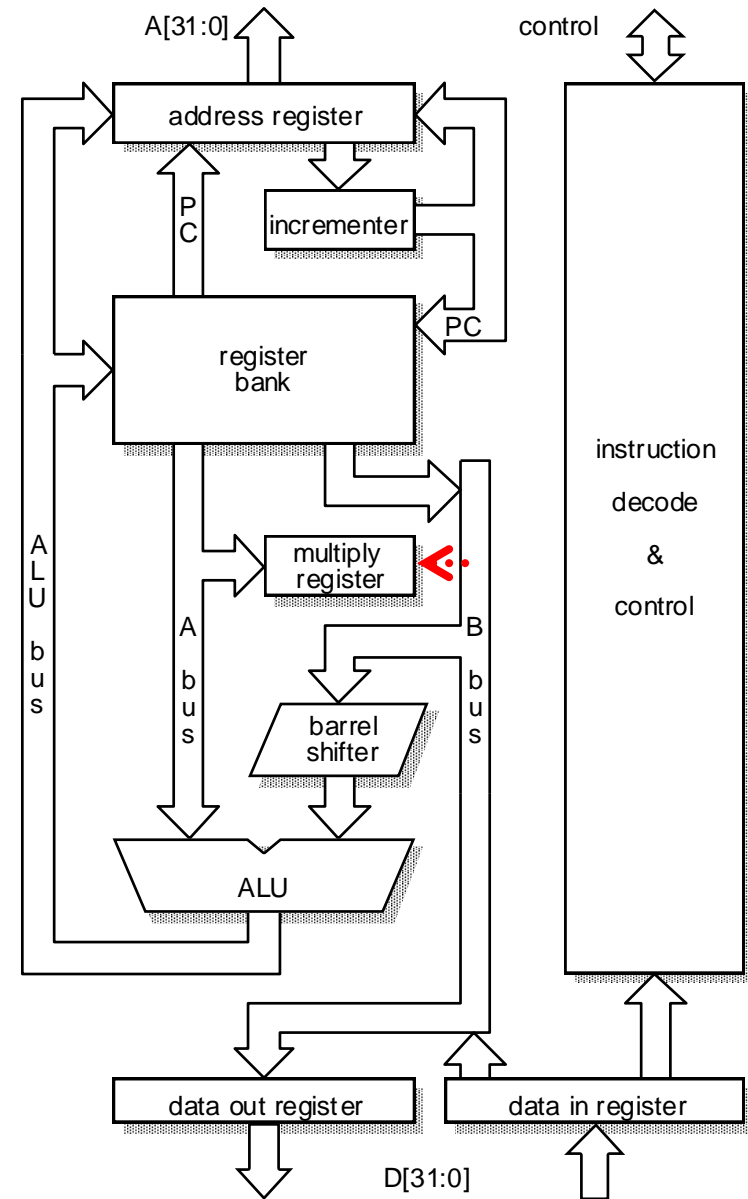


ARM7 CPU:

many internal registers;
address register
and register bank.

3 internal datapaths;
A bus,
B bus,
ALU bus.

additional ALU functionality;
multiplier
and barrel shifter.



The barrel shifter

- The barrel shifter is a very useful feature of the ARM7 microprocessor which allows bit patterns to be rotated.
- E.g. the instruction

MOV r1, r2, LSL #5

- would take the bit pattern in register r2 and shift it 5 places to the left before placing it in register r1. So if r2 held:

1010 0101 1100 0011 1001 0110 1110 0111

- after the instruction was executed r1 would hold:

1011 1000 0111 0010 1101 1100 1110 0000

ALU adder design

- ripple-carry adder
- carry look-ahead
- carry select adder

Booth's algorithm: Nth cycle

Booth's algorithm works by replacing $\times 3$ by $\times (4-1)$ so that each cycle multiplies by a 2 bit value.

N bit multiplications can be completed in $N/2$ cycles.

The table right summarizes the action on the Nth cycle.

Table 4.3 in Furber also replaces $\times 2$ by $\times (4-2)$ for $C_{in}=0$ as this makes the control logic slightly simpler.

C_{in}	Multiplier	LSL #	ALU	C_{out}
0	$\times 00_2$	-	$A+0$	0
0	$\times 01_2$	$2N$	$A+B$	0
0	$\times 10_2$	$(2N+1)$	$A+B$	0
0	$\times 11_2$	$2N$	$A-B$	1
1	$\times 00_2$	$2N$	$A+B$	0
1	$\times 01_2$	$(2N+1)$	$A+B$	0
1	$\times 10_2$	$2N$	$A-B$	1
1	$\times 11_2$	-	$A+0$	1

- ✓ Organization/Architecture
- ✓ **Mnemonics**
- ✓ Flags
- ✓ Representations
- ✓ Subroutines
- ✓ Interrupt
- ✓ Pipeline
- ✓ Encoding
- ✓ Memory
- ✓ Processor Modes

Program counter.

- Instructions are stored in memory so:
 - How does the CPU know the memory address for the next instruction in the computer program?
 - Register r15 always holds the memory address of the next instruction to be executed.
 - An alternative name for register r15 is the 'program counter'.

Instructions

- For a human to perform a task he/she needs instructions e.g.
 - to knit a jumper
 - the instructions are the knitting pattern
 - to bake a cake
 - the instructions are in the recipe.
- Computers also need instructions
 - so that the data is processed into information correctly.

Instructions stored in memory.

- Instructions are 32 bits long, whereas memory locations are 8 bits long so one instruction occupies four locations in memory, e.g.
 - one instruction would be stored in 4 memory locations with addresses 0x00008000, 0x00008001, 0x00008002 and 0x00008003.
 - The next instruction would be stored at addresses 0x00008004, 0x00008005, 0x00008006 and 0x00008007 and so on.

Simple instructions.

- One of the simplest instructions is to move a value into a register e.g.
 - move 114 into register r12
 - The machine code for this instruction is 0xE3A0C072
 - After the instruction is executed
 - register r12 will hold the value 0x00000072 (hexadecimal for 114) and
 - value in register r15, the program counter, will have increased by 4.
 - All other registers remain unchanged.

Machine code.

- Look at the machine code for the instruction;
 - move 114 into register r12, again.
 - 0xE3A0C072
 - 1110 0011 1010 0000 1100 0000 0111 0010
 - The value 114 is given in the least significant byte
 - 0xE3A0C072
 - 114 in decimal is 72 in hexadecimal.
 - Register r12 is given by the 5th digit
 - 0xE3A0C072
 - 12 in decimal is C in hexadecimal.

Mnemonics

- In general nobody remembers all of the machine code for any particular processor (or indeed any).
- Instead we use mnemonics
 - mnemonics are words or phrases which are easy to remember and
 - can replace something which is difficult to remember.
- The instruction to move the value 114 into register r12 has the mnemonic
 - MOV r12, #114.
 - This is much easier to remember than 0xE3A0C072

Instructions for Arithmetic

- The ARM7 can add, subtract and multiply numbers (but not divide).
- The mnemonic for add the value in register **x** to the value in register **y** and place the sum in register **z** is:
ADD **r****z**, **r****y**, **r****x**
- E.g. to add the value in register r1 to the value in register r2 and leave the sum in register r3
 - the mnemonic is: ADD r3, r2, r1
 - We don't need to know the machine code.

Instructions using logic

- As well as arithmetic the ARM7 can also do logic such as AND and OR.
- AND value in r x with value in r y and leave the result in r z has the mnemonic:

AND r z , r y , r x

- OR value in r x with value in r y and leave the result in r z has the mnemonic:

ORR r z , r y , r x

- ✓ Organization/Architecture
- ✓ Mnemonics
- ✓ **Flags**
- ✓ Representations
- ✓ Subroutines
- ✓ Interrupt
- ✓ Pipeline
- ✓ Encoding
- ✓ Memory
- ✓ Processor Modes

Flags

Flags are used to give a signal that something either has or has not happened.

For example, the flag is raised to half mast on in front of buildings as a signal that someone important is dead.

Microprocessors also use “flags” to signify what has happened.

Flags and Processors

Flags are essentially 1 bit memory devices and different microprocessors have different flags.

The ARM microprocessor has four flags that are commonly found in the majority of microprocessors.

These flags are the zero flag (Z), the negative flag (N), the carry flag (C), and the overflow flag (V).

Flags are also known as 'condition codes'.

What are flags used for?

Flags are used in two ways:

The main use of flags is to determine if another instruction is executed or not. This is called conditional execution.

The carry flag can also be used in some arithmetic instructions as an additional value (we will return to this later).

Branches

- An important use of conditional execution is in branches.
- Conditional branches are used when a processor performs one function rather than another
 - e.g.
 - display text messages rather than receive an incoming call.
- First consider the unconditional branch.

Flags - summary.

- The zero flag, Z, is set when the result (not including any carry out) is 0x00000000.
- The negative flag, N, is set when the most significant bit of the 32 bit result is 1.
- The carry flag, C, is set when the result (taken as an unsigned integer) is greater than $(2^{32} - 1)$.
- The overflow flag, V, is set when the result (taken as a two's complement number) is greater than $(2^{31} - 1)$ or less than -2^{31} .

- ✓ Organization/Architecture
- ✓ Mnemonics
- ✓ Flags
- ✓ **Representations**
- ✓ Subroutines
- ✓ Interrupt
- ✓ Pipeline
- ✓ Encoding
- ✓ Memory
- ✓ Processor Modes

Representing characters

- As well as manipulating numbers, computers also manipulate characters e.g.
 - in a word processing package like Microsoft Word.
 - Characters must be coded in binary before computers can process them
 - the standard coding for characters is ASCII (American Standard Code for Information Interchange).

Negative numbers

- There are two main methods for representing negative numbers in microprocessors.
- These are:
 - 1) Sign magnitude.
 - 2) Two's complement.
- In each case the most significant bit - 'm.s.b.' - indicates the sign (1 for -ve and 0 for +ve).

2's complement

- The two's complement method automatically sets the m.s.b. or 'sign bit' to 1.
- The following method can be used to find a 2's complement representation of a negative number,
 - e.g.
 - -20640
 - First find the positive value: 0x000050A0 or 0000 0000 0000 0000 0101 0000 1010 0000₂
 - Next invert all bits (0 → 1, 1 → 0).
 - 1111 1111 1111 1111 1010 1111 0101 1111₂ or
 - 0xFFFFAF5F.
 - And then add 1 → 0xFFFFAF60

Floating point numbers

- The floating point format uses the equation $A \times 2^n$ to represent numbers.
- Using the IEEE 754 standard the A value has a sign and 24 bits and the exponent value, n, has 8 bits.
- The number $5,637,144,576_{10}$ ($= 2^{32} + 2^{30} + 2^{28}$) is too big for a 32 bit unsigned integer.
- In binary this number is
 - $1\ 0101\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000_2$
- First we must normalize this number as follows:

5,637,144,57610 in IEEE 754

- $1\ 0101\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000_2 = 1.0101\ 0000\ 0000\ 0000\ 0000\ 000_2 \times 2^{32}$
- The exponent is $32_{10} = 100000_2$ but in the IEEE 754 format we must add $127_{10} = 1111111_2$ to this to give $10011111_2 (=159_{10})$
- In 32 bits this becomes:
 $0100\ 1111\ 1010\ 1000\ 0000\ 0000\ 0000\ 0000$
- Note that the normalization means that the m.s.b. of the A value is always 1 so this can be omitted.
- The m.s.b. of the 32 bit word is a sign bit.

- ✓ Organization/Architecture
- ✓ Mnemonics
- ✓ Flags
- ✓ Representations
- ✓ **Subroutines**
- ✓ Interrupt
- ✓ Pipeline
- ✓ Encoding
- ✓ Memory
- ✓ Processor Modes

Subroutines

- Many microprocessor programs include groups of instructions which are repeated many times.
- It is wasteful of memory to include these instructions in the program again and again.
- Instead they can be included once in a special structure known as a **subroutine**.
- The main program branches to the start of the subroutine when it requires those particular instructions.
- At the end of the subroutine there is another branch back to the main program.

Problem

- At the end of the subroutine how does the processor know which instruction to return to?
- This problem is solved by using a 'link register'.
- A link register holds the memory address of the instruction in the main program to which the subroutine returns to.
- Register **r14** in the ARM microprocessor is designated as the link register and 'r14' can be replaced by '**lr**' in mnemonics.

Branch and link

- The mnemonic for ‘branch and link’ is BL;
 - a simple branch with mnemonic B does not update the link register.
- To return from the subroutine the value in the link register is moved into the program counter;
 - MOV pc, lr.
- What happens if a branch and link occurs in a subroutine?
 - The value held in the link register will be overwritten by a new return address so before one subroutine calls another the link register value must be stored elsewhere.

Nested subroutines

- Complicated programs will have several subroutines and some subroutines will call other subroutines
 - this is standard practice and it is known as nesting.
- If subroutine A calls subroutine B the link register can not store the return address for both subroutines at the same time.
- In order to preserve the return address of all subroutines an area of computer memory called the **stack** is used.
- The stack is a '**last in first out**' queue.

The stack

- The stack is a last in first out queue
 - that means whatever data was added to the stack ('pushed') last is taken from the stack ('popped') first.
 - E.g. if the values pushed onto a stack were 0x00FF, 0xFF00, 0xAAAA in that order then they would be popped from the stack in the reverse order.
- In memory the stack is held as a list:

top of stack 0xAAAA

0xFF00

bottom of stack 0x00FF

Another problem

- The stack is a dynamic memory structure
 - Meaning that the amount of data held in the stack can vary.
- The bottom of the stack can be fixed at a particular memory address.
- How do we know where the top of the stack is?
 - We need to know so that we can pop the correct data.
- Another register is used to identify the top of the stack - it is known as the **stack pointer**.

- ✓ Organization/Architecture
- ✓ Mnemonics
- ✓ Flags
- ✓ Representations
- ✓ Subroutines
- ✓ **Interrupt**
- ✓ Pipeline
- ✓ Encoding
- ✓ Memory
- ✓ Processor Modes

Interrupts

- Many applications for microprocessor systems require things to happen when an event occurs unexpectedly.
- These unexpected events can interrupt the normal operation of the microprocessor.
- During an 'interrupt' the execution of the main program is halted and a special program is executed instead.
- When the interrupt program has finished, the microprocessor returns to the main program.

Interrupts

- Interrupts provide a very convenient method for dealing with events and this method is often used for events that are not unexpected
 - e.g.
 - when a mobile phone receives an incoming call.
- An interrupt is triggered when the microprocessor receives a (voltage) signal on a special connection within the control bus.
- The ARM7 has two types of interrupts,
 - a normal interrupt (IRQ) and
 - a fast interrupt (FIQ).

Interrupt Handling

- When an interrupt occurs the following happens:
 - ① The registers for IRQ mode or FIQ mode are activated.
 - ② The current program status register, CPSR, (this contains the flags and other information) is saved into a saved program status register, SPSR.
 - a) There are two SPSR registers, one for each mode.
 - ③ The return address of the next instruction to be executed in the main program is stored in the link register for the appropriate mode.
 - ④ The program counter is set to either 0x00000018 for an IRQ or 0x0000001C for a FIQ.

Interrupt Vectors

- The first instruction executed by an interrupt program is the instruction stored at memory address 0x00000018 for an IRQ or 0x0000001C for a FIQ.
- These memory addresses are known as 'vectors'.
- For an IRQ the instruction at address 0x00000018 must be a branch to another part of memory because the following memory location, 0x0000001C contains the first instruction of the FIQ handler.

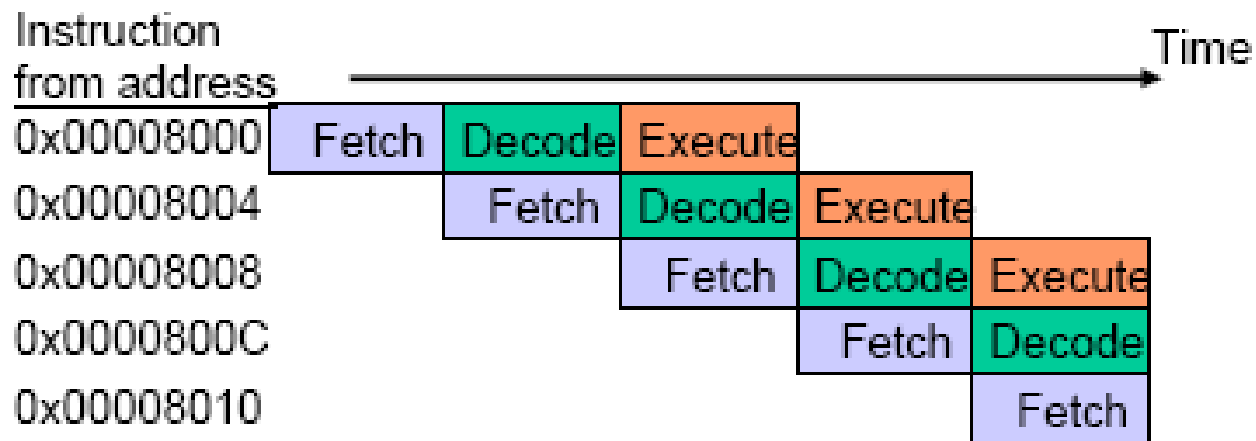
- ✓ Organization/Architecture
- ✓ Mnemonics
- ✓ Flags
- ✓ Representations
- ✓ Subroutines
- ✓ Interrupt
- ✓ **Pipeline**
- ✓ Encoding
- ✓ Memory
- ✓ Processor Modes

Instruction Pipelines

- Instruction pipelines are an important feature of all modern microprocessors.
- For the same basic speed of transistor operation, an **n stage instruction pipeline** allows the microprocessor to execute up to **n times as many instructions** in a given time.
- The ARM7 microprocessor has a three stage pipeline
 - one stage for each of the CPU cycles;
 - fetch, decode and execute.

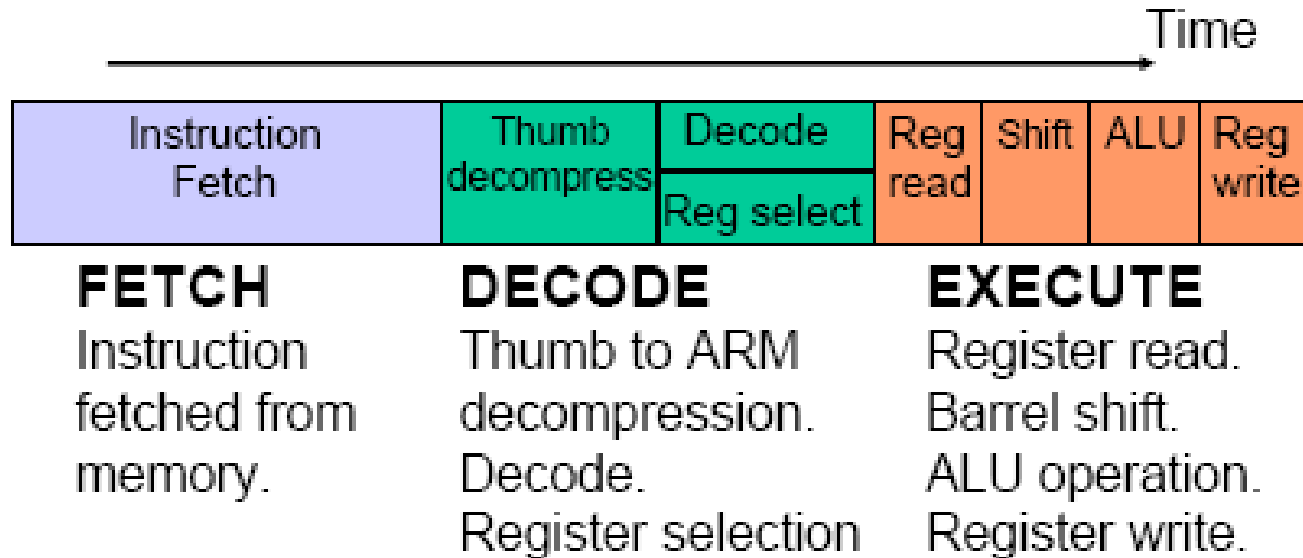
Instruction Pipelines

- In a three stage pipeline, the CPU can simultaneously execute an instruction, decode the next instruction and fetch the next instruction.



ARM7 3 stage pipeline: detail

In each stage of the ARM7 pipeline, several things happen; normally consecutively:



- ✓ Organization/Architecture
- ✓ Mnemonics
- ✓ Flags
- ✓ Representations
- ✓ Subroutines
- ✓ Interrupt
- ✓ Pipeline
- ✓ **Encoding**
- ✓ Memory
- ✓ Processor Modes

ARM7 pipeline example

Consider a typical three operand ARM7 instruction e.g.

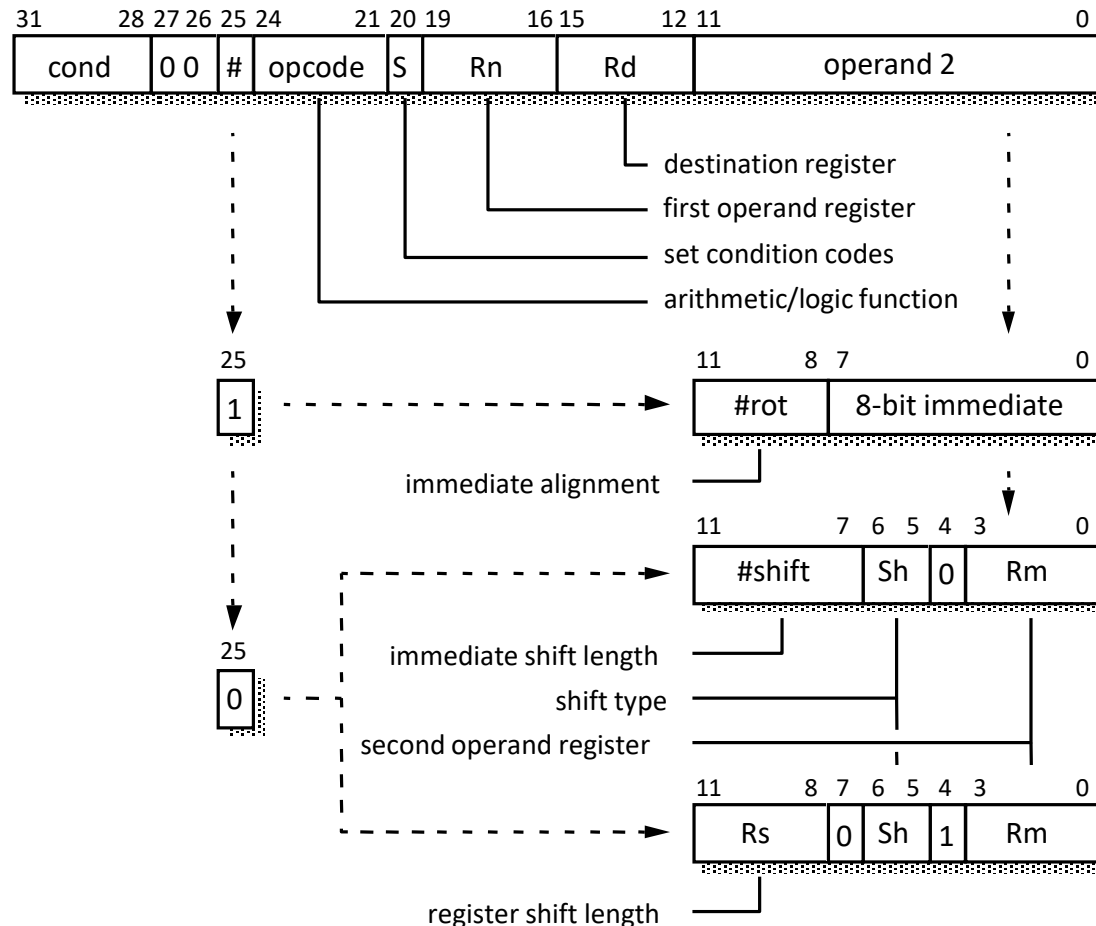
ADDNE r11, r3, r7 LSR #8

Add the contents of register r3 to the contents of register r7 shifted right by 8 bits and put the sum in register r11 if the zero flag is clear (Z=0).

The machine code for this instruction is 0x1083B427



Data processing instruction binary encoding



ARM data processing instructions

Opcode [24:21]	Mnemonic	Meaning	Effect
0000	AND	Logical bit-wise AND	$Rd := Rn \text{ AND } Op2$
0001	EOR	Logical bit-wise exclusive OR	$Rd := Rn \text{ EOR } Op2$
0010	SUB	Subtract	$Rd := Rn - Op2$
0011	RSB	Reverse subtract	$Rd := Op2 - Rn$
0100	ADD	Add	$Rd := Rn + Op2$
0101	ADC	Add with carry	$Rd := Rn + Op2 + C$
0110	SBC	Subtract with carry	$Rd := Rn - Op2 + C - 1$
0111	RSC	Reverse subtract with carry	$Rd := Op2 - Rn + C - 1$
1000	TST	Test	Scc on $Rn \text{ AND } Op2$
1001	TEQ	Test equivalence	Scc on $Rn \text{ EOR } Op2$
1010	CMP	Compare	Scc on $Rn - Op2$
1011	CMN	Compare negated	Scc on $Rn + Op2$
1100	ORR	Logical bit-wise OR	$Rd := Rn \text{ OR } Op2$
1101	MOV	Move	$Rd := Op2$
1110	BIC	Bit clear	$Rd := Rn \text{ AND NOT } Op2$
1111	MVN	Move negated	$Rd := \text{NOT } Op2$

- ✓ Organization/Architecture
- ✓ Mnemonics
- ✓ Flags
- ✓ Representations
- ✓ Subroutines
- ✓ Interrupt
- ✓ Pipeline
- ✓ Encoding
- ✓ **Memory**
- ✓ Processor Modes

What is computer memory?

- Computer memory is a very big sequential logic circuit made up of thousands or millions of simple logic gates, such as a D type latch, which can remember a 0 or a 1, that is one bit of data.
- Groups of these gates are collected together in a memory 'location'.
- There are typically 8 bits of data in one location.
- Each memory location has a unique memory address.

Memory organization.

- Taking the ARM7TDMI microprocessor as an example
 - at each memory location
 - it has 8 bits of data
 - 8 bits is known as a byte.
- The ARM is a 32 bit processor and addresses are 32 bits long from 0x00000000 to 0xFFFFFFFF.
 - That means there can be up to 4,294,967,296 (or 2^{32}) different memory locations all with a unique memory address.
- In practice not all addresses are used for memory.

Addressing modes

- The concept of addressing mode is common to all microprocessors and an understanding of addressing mode is important for users of microprocessors.
- The ARM7 microprocessor supports many different addressing modes and we will concentrate on just four.

Load and Store

- When a load instruction is executed
 - the data travels from memory to register
- whereas when a store instruction is executed
 - the data travels from register to memory.
- Each memory location holds one byte or 8 bits of data whereas each register holds 4 bytes of data.
- So LDR and STR use four consecutive memory addresses but which byte goes to which location?

Little endian

- Microprocessors can be either 'little endian' or 'big endian'
- If the processor is 'little endian' then the instruction:
 - STR r6, [r11]
 - with 0xFFAABB11 in r6 and 0x00008000 in r11
 - would store
 - byte 0x11 at address 0x00008000
 - 0xBB at 0x00008001
 - 0xAA at 0x00008002
 - 0xFF at 0x00008003

Big endian

- Whereas if the processor is 'big endian' then the instruction:
 - STR r6, [r11]
 - with 0xFFAABB11 in r6 and 0x00008000 in r11
 - would store
 - byte 0xFF at address 0x00008000
 - 0xAA at 0x00008001
 - 0xBB at 0x00008002
 - 0x11 at 0x00008003

Types of Silicon Memory

- Silicon IC memory can be classified as either read-only or read-write.
- Read-only memory (or ROM) is used for the operating system on a desktop computer or an application program in an embedded system such as a mobile phone.
- Read-write memory (which is normally called RAM) is used for an application program on a desktop computer and for temporary data.

Types of ROM

- There are several generations of silicon ROM.
- Programmable ROM (PROM) could be electrically programmed **once only** and then the contents were fixed.
- Erasable PROM (EPROM) could also have its contents erased by exposure to **ultraviolet** (UV) light and could then be **reprogrammed**.
- Electrically erasable PROM (EEPROM) could have its contents erased by an **electrical signal**.

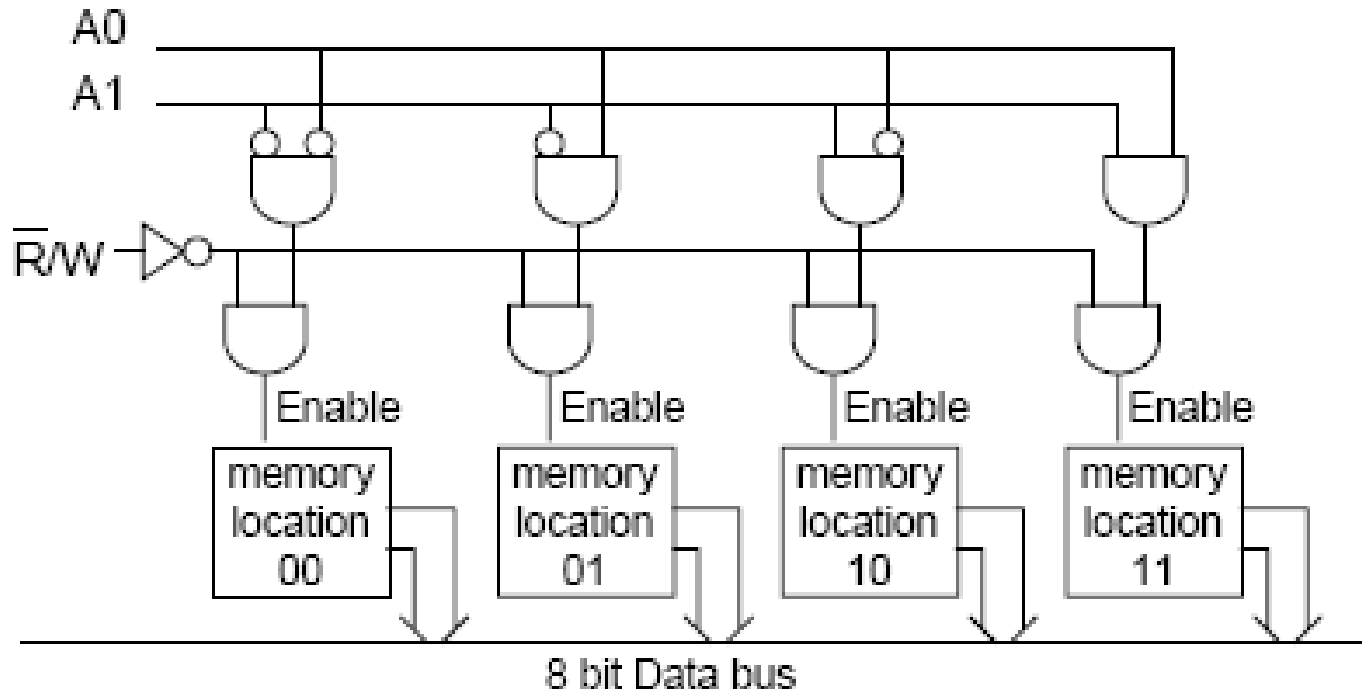
Types of ROM

- Flash Erasable PROM (FEPROM or flash memory) is **similar to EEPROM** but erasing could only be done over a significant part of (maybe all) the integrated circuit.
- All ROM is non-volatile - meaning the information stored in ROM memory is not lost when it is disconnected from a power source. The data stored in ROM will remain almost indefinitely e.g. many tens of years.
- In contrast the contents of RAM memory will be lost almost as soon as power is switched off - RAM is volatile.

Cache memory

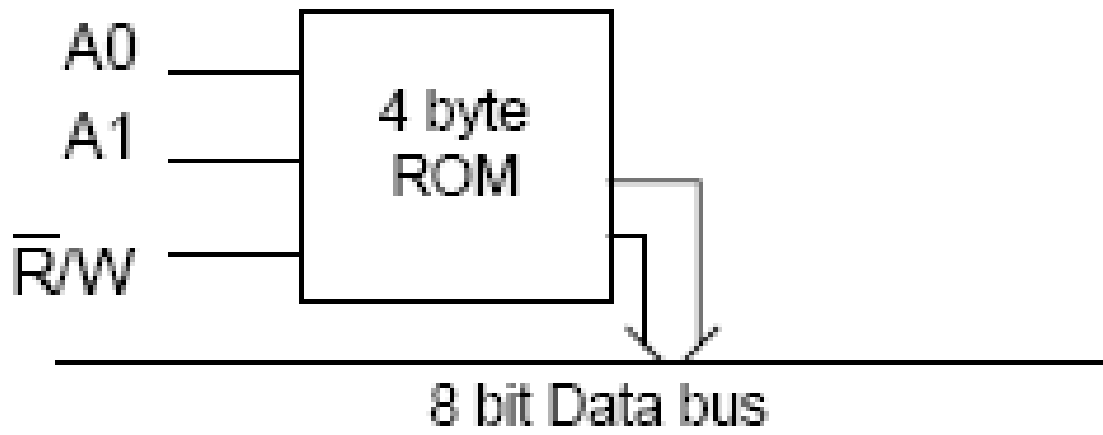
- When data is read from, or written to, main memory a copy of the data is saved in the memory cache (along with the main memory address).
- When a subsequent read occurs the cache can be checked to see if the required data is already there.
- If it is (a cache 'hit') then it can be supplied immediately.
- If not (a cache 'miss') then it must be fetched from main memory.

Simple 4 byte ROM addressing



Simple ROM addressing

- So memory location 00 is 'enabled' and connected to the data bus only when $R^*/W = 0$, $A0 = 0$ and $A1 = 0$.
- The previous circuit can be represented by a single block as follows:



Tri-state gates

- Only one memory location should be connected to the bus at one time – one device can not drive a logic 0 onto a bus line at the same time as another device drives a logic 1 onto the same line.
- This is achieved using ‘tri-state gates’ which have an ‘enable’ input – when not enabled the output acts like an ‘open circuit’ which is referred to as ‘high impedance’.

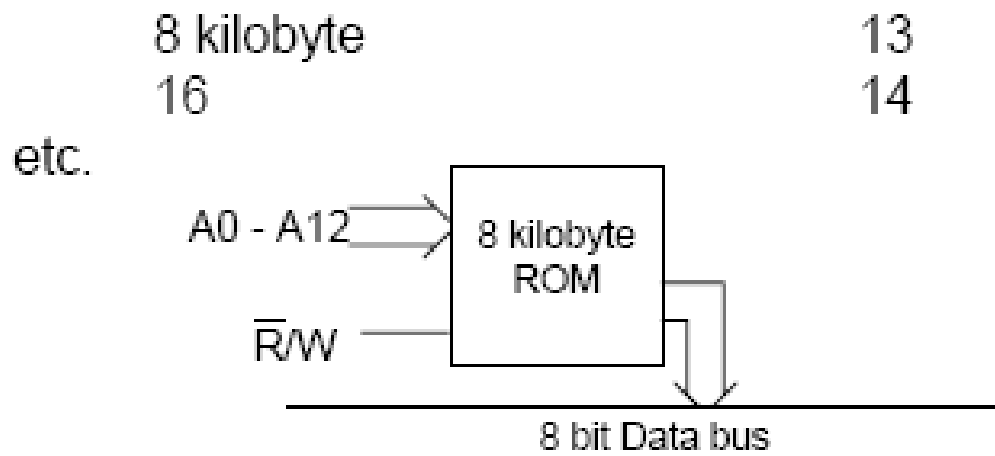
Tri-state gates



enable	input	output
0	0	high-impedance
0	1	high-impedance
1	0	0
1	1	1

Simple ROM addressing

- Bigger memories would have similar connections but there would be more address lines.
- A 2^N byte memory would need N address lines:
- E.g. a 4 kilobyte memory would need 12 address lines.



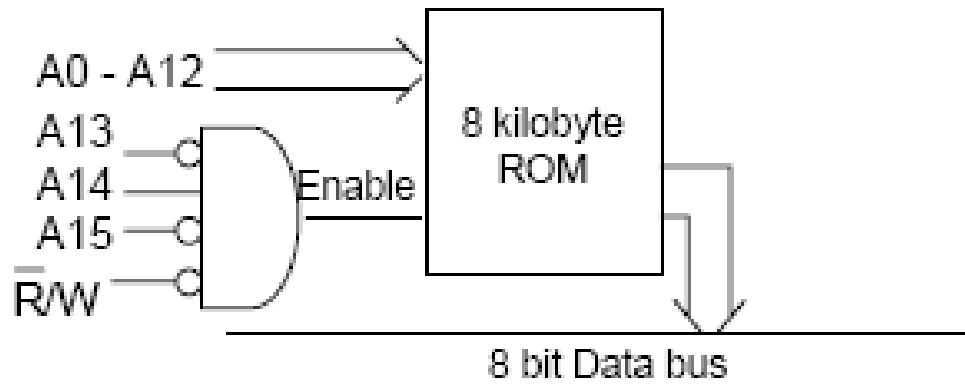
Which address?

- If an 8 kilobyte memory chip is connected to a processor with a 16 bit address bus whereabouts is the memory located within in the 64 kilobyte address space?
- The lowest N address lines are connected to the 2^N byte memory; in this case A0 to A12. The other address lines, A13 to A15, must be decoded e.g. if the memory range is from 0x4000 to 0x5FFF or in binary:


A15	A14	A13	A12	A11	A10	A9	A8	A7	A6	A5	A4	A3	A2	A1	A0
0	1	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	1	0	1	1	1	1	1	1	1	1	1	1	1	1	1

8 kilobyte ROM addressing

- The memory is only enabled when $A13 = 0$, $A14 = 1$, $A15 = 0$ and $R^*/W = 0$.




- If the address bus is 32 bits wide then all the address lines not connected to the memory chip, A13 to A31, are decoded.



Connecting to a 32 bit data bus

- Memory with 8 data bits output can be connected to a 32 bit data bus by using 4 separate memory chips; one for each byte



$A_0 - A_{n+1}$: memory size

$A_{n+2} - A_{31}$: location in the memory map

A_0
 A_1

Memory maps

- The location of a memory chip within the complete addressable space can be represented by a memory map.
- E.g. for the 8 kilobyte ROM memory between addresses 0x4000 and 0x5FFF in a 64 kilobyte addressable space:

0xFFFF	empty
0x6000	
0x5FFF	8k ROM
0x4000	
0x3FFF	empty
0x0000	

- ✓ Organization/Architecture
- ✓ Mnemonics
- ✓ Flags
- ✓ Representations
- ✓ Subroutines
- ✓ Interrupt
- ✓ Pipeline
- ✓ Encoding
- ✓ Memory
- ✓ **Processor Modes**

Different registers in each mode

User mode registers

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)
r15 (pc)

FIQ
mode

r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)

IRQ
mode

r13 (sp)
r14 (lr)

SVC
mode

r13 (sp)
r14 (lr)

Undef
mode

r13 (sp)
r14 (lr)

Abort
mode

r13 (sp)
r14 (lr)

cpsr

spsr

spsr

spsr

spsr

spsr

System mode uses
the same registers
as user mode.

in user mode

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
r13 (sp)
r14 (lr)
r15 (pc)

cpsr